




This course was developed as a part of the VLIR-UOS Cross-Cutting projects:

- Statistics: 2011-2016, 2017.
- Statistics: 2017.
- Statistics for development : 2018-2020.



The >eR-Biostat initiative

Making R based education materials in
statistics accessible for all

Applied Generalized Linear Models (GLM) using R (PART 2)

Developed by

Tadesse Awoke (Gondar University), Said Mussa (Mekele University) , Ziv Shkedy
(Hasselt University), and Fetene Tekle (J & J)

LAST UPDATE: 07/09/2022



ER-BioStat

Email: erbiostat@gmail.com



<https://github.com/eR-Biostat>



@erbiostat

Reference list

- Main reference
 - Dobson (2002): An introduction to generalized linear models.
- Other references:
 - McCullagh and Nelder (1983): Generalized linear models (first edition).
 - Collet D(1994): Modeling Binary data.
 - Lindsey (1997): Applying generalized linear models.



Software

- Two main R functions:
 - Linear models in R: the `lm()` function.
 - Generalized linear models in R: the `glm()` function in R.
- All R programs for the examples presented in the slides are available online.



Datasets

- Data are given as a part of R programs for the course.
- External datasets (which are not given as a part of the R code) and used for illustration are available online.

Topics (part 2)

- 11. Poisson Regression
- 12. Beyond Poisson and binomial distributions: models with different link functions and/or distributions
- 13. Poisson regression and log linear models
- 14. Over dispersion

Chapter 11: Poisson Regression

Donson: chapter 7.

Lindsey: Appendix B.

McCullagh & Nelder: chapter 2.

Count data

Count data:

- counts per unit of time/area/distance, etc
- contingency tables: counts cross-classified by categorical variables
- Covariates: categorical or continuous


Example 1: Stress data

- One randomly chosen member from each randomly chosen household in a sample from Oakland, California, USA was interviewed. In a list of 41 events, respondents were asked to note which had occurred within the last 18 months. The result is given as:

Month	1	2	3	4	5	6	7	8	9
Respondents	15	11	14	17	5	11	10	4	8
Month	10	11	12	13	14	15	16	17	18
Respondents	10	7	9	11	3	6	1	1	14

Data in R

```
> stress <- read.table("C:...../stress.txt", sep=";", header=TRUE)
> attach(stress)
plot(respondents ~ month, xlab = "Months", ylab = "Subjects",
     xlim=c(0,20), ylim=c(0,20))
```



Model formulation

The distribution of the response variable

$$Y_t \sim \text{Poisson}(\mu_t)$$

$$f(Y_t, \mu) = \frac{\mu^{y_i} e^{-\mu}}{Y_t!}$$

$$E(Y_t) = \mu_t$$

The dependency on the predictor

$$\mu_t = f(\text{month})$$

A proposal ?

$$\mu_t = \beta_0 + \beta_1 t$$

Model formulation

The distribution of the response variable

$$Y_t \sim \text{Poisson}(\mu_t)$$

$$E(Y_t) = \mu_t$$

The linear predictor

$$\mu_t = f(\text{month})$$

$$\eta = \beta_0 + \beta_1 M_t$$

$$\mu_t = e^\eta = e^{\beta_0 + \beta_1 M_t}$$

$$g(E(Y_t)) = \log(\mu_t) = \eta$$

GLM for Poisson regression using glm()

$$\eta = \beta_0 + \beta_1 t$$

> respGLM <- glm(respondents ~ month,
family=poisson, data=stress)

The relative risk

$$RR = \frac{E(Y_t | t+1)}{E(Y_t | t)} = \frac{e^{\beta_0 + \beta_1(t+1)}}{e^{\beta_0 + \beta_1 t}} = e^{\beta_1}$$

GLM for Poisson regression using glm()

```
> summary(respGLM)
```

Call:

```
glm(formula = respondents ~ month, family = poisson, data = stress)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.9886	-0.9631	0.1737	0.5131	2.0362

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	2.80316	0.14816	18.920	< 2e-16 ***
month	-0.08377	0.01680	-4.986	6.15e-07 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 50.843 on 17 degrees of freedom


Residual deviance: 24.570 on 16 degrees of freedom

AIC: 95.825

Number of Fisher Scoring iterations: 5

Data and fitted model

```
plot(respondents ~ month, xlab = "Months",  
     ylab = "Subjects", xlim=c(0,20), ylim=c(0,20))  
lines(month,respGLM$fit)
```



Example 2: Ministerial resignation

- On October 18, 1995, 'The Independent' reported on the numbers of ministerial resignations because of different reason. The years start in 1945-1951, with a Labour government, and 7 Resignations.

Term	45-51	51-57	55-57	57-63	63-64	64-70	70-74	74-76	76-79	79-90	90-95	97-05
Gov	Lab	con	con	con	con	lab	con	lab	lab	con	con	lab
Res	7	1	2	7	1	5	6	5	4	1	1	1
Year	6	4	2	6	1	6	4	2	3	1	5	8

- Main question: Is there any difference between Government (Labor and Conservative) in the rate of resignations?

Model formulation

The distribution of the response variable

$$Y_t \sim \text{Poisson}(\mu_t)$$

$$E(Y_t) = \mu_t$$

The linear predictor

$$\mu_t = f(\text{gov}: L / C)$$

$$G_t = \begin{cases} 1 & L \\ 0 & C \end{cases}$$


$$\eta = \beta_0 + \beta_1 G_t$$

$$\mu_t = e^\eta = e^{\beta_0 + \beta_1 G_t}$$

$$g(E(Y_t)) = \log(\mu_t) = \eta$$


The data in R

```
>Resignations <- read.table("C:...../resign.txt",header=T)  
>attach(Resignations)  
>plot(Res ~ log(years), pch=19, col=c(4,2)[Gov])
```



Model formulation (1)

```
> first.glm <- glm(Res ~ Gov , poisson)
```


$$\eta = \log(\mu_t) = \beta_0 + \beta_1 G_t$$

The relative risk

$$RR = \frac{E(Y_t | L)}{E(Y_t | C)} = \frac{e^{\beta_0 + \beta_1}}{e^{\beta_0}} = e^{\beta_1}$$

Model 1 output in R

```
> first.glm <- glm(Res ~ Gov, poisson);
```

```
> summary(first.glm)
```

Call:

```
glm(formula = Res ~ Gov, family = poisson)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.5331	-1.2942	-0.3255	0.7548	2.7793

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.79176	0.15430	11.61	<2e-16 ***
Govlab	0.09531	0.23262	0.41	0.682

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 33.436 on 11 degrees of freedom

Residual deviance: 33.269 on 10 degrees of freedom

AIC: 78.459

Number of Fisher Scoring iterations: 5

Model formulation (2)

How can we take the number of government years into account ?

$$\mu_t = e^\eta = e^{\beta_0 + \beta_1 G_t + \beta_2 \log(\text{years}_t)}$$

$$\eta = \log(\mu_t) = \beta_0 + \beta_1 G_t + \beta_2 \log(\text{years}_t)$$

GLM with Poisson family

```
> first.glm <- glm(Res ~ Gov + log(years), poisson)
```



The same slope for log(year)

$$\eta = \log(\mu_t) = \beta_0 + \beta_1 G_t + \beta_2 \log(years_t)$$

Model 2: R output

```
glm(formula = Res ~ Gov + log(years), family = poisson)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.2417	-0.3469	-0.1250	0.3917	1.6513

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	0.29973	0.41598	0.721	0.471
Govlab	0.03541	0.23271	0.152	0.879
log(years)	0.96636	0.22258	4.342	1.41e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 33.436 on 11 degrees of freedom
Residual deviance: 11.276 on 9 degrees of freedom
AIC: 58.466

Number of Fisher Scoring iterations: 4

Model 3: model formulation

Different intercepts and slopes

```
> first.glm3 <- glm(Res ~ log(years)+Gov+Gov:log(years),  
  poisson)
```

$$\eta = \beta_0 + \beta_1 G_t + \beta_2 \log(\text{years}_t) + \beta_3 G_t \log(\text{year}_t)$$

Model 3 output in R

```
> summary(first.glm)
```

Call:

```
glm(formula = Res ~ log(years) + Gov + Gov:log(years), family = poisson)
```

Estimate Std. Error z value Pr(>|z|)

(Intercept) 0.04443 0.51085 0.087 0.931

log(years) 1.11144 0.27306 4.070 4.69e-05 ***

Govlab 0.81973 0.82744 0.991 0.322

log(years):Govlab -0.46049 0.46880 -0.982 0.326

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 33.436 on 11 degrees of freedom

Residual deviance: 10.336 on 8 degrees of freedom

AIC: 59.526

Number of Fisher Scoring iterations: 4

Model 4: GLM with an offset variable: model formulation

Number of resignation per government year:

$$\frac{Y_t}{year_t} \sim \text{Poisson}(\mu_t)$$

$$Y_t \sim \text{Poisson}(years_t \times \mu_t)$$


Model 4: GLM with an offset variable - model formulation

$$Y_t \sim \text{Poisson}(\text{years}_t \times \mu_t)$$

$$E(Y_t) = \text{years}_t \times \mu_t = \text{years}_t \times e^{\beta_0 + \beta_1 G_t}$$

$$g(E(Y_t)) = g(\text{years}_t \times \mu_t) = \log(\text{years}_t \times e^{\beta_0 + \beta_1 G_t})$$

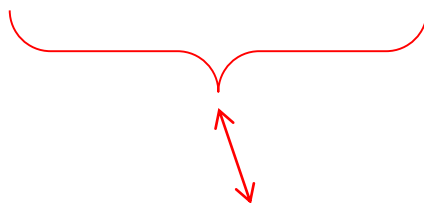
$$g(E(Y_t)) = \log(\text{years}_t) + \beta_0 + \beta_1 G_t = \eta$$



offset variable

A model with offset in R

```
> next.glm<- glm(Res ~ Gov + offset(log(years)), poisson)
```


$$\beta_0 + \beta_1 G_t + \log(years_t)$$

Model 4 GLM with offset output in R

```
summary(first.glm4)
```

Call:

```
glm(formula = Res ~ Gov + offset(log(years)), family = poisson)
```

	Estimate	Std. Error	z value	Pr(> z)
$\hat{\beta}_0$ → (Intercept)	0.24116	0.15430	1.563	0.118
$\hat{\beta}_1$ → Govlab	0.03647	0.23262	0.157	0.875

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 11.323 on 11 degrees of freedom

Residual deviance: 11.299 on 10 degrees of freedom


AIC: 56.488

Number of Fisher Scoring iterations: 4

Data and predicted models

```
plot(Res ~ log(years), pch=19, col=c(4,2)[Gov]) # Use palette() to find out which colour corresponds  
> legend(locator(1), legend= c("conservative", "labour"), col=c(4,2), pch=19)  
> l <- (0:25)/10  
> fv <- exp(0.3168 + 0.9654*l)# to plot fitted curve under last.glm  
> lines(l,fv)
```

Ministerial Resignations against log(years)



AIC for the different models

MODEL	No parameters	Deviance	AIC
1	2	33.269	78.45861
2	3	11.276	58.46574
3	4	10.336	59.52603
4	2	11.299	56.48846


$$g(E(Y_t)) = \log(\text{years}_t) + \beta_0 + \beta_1 G_t$$

Confidence interval for β_2 for model 2

```
first.glm2 <- glm(Res ~ log(years)+Gov, poisson); summary(first.glm)
```

```
> confint(first.glm2, level=0.95)
```

```
Waiting for profiling to be done...
```

```
2.5 % 97.5 %
```

```
(Intercept) -0.5701863 1.0636716
```

```
log(years) 0.5447242 1.4185361
```

```
Govlab -0.4268576 0.4893354
```


ANOVA for model 3 and model 2

```
> anova(first.glm3,first.glm2, test = "Chisq")
```

Analysis of Deviance Table

Model 1: Res ~ log(years) + Gov + Gov:log(years)

Model 2: Res ~ log(years) + Gov

Resid. Df Resid. Dev Df Deviance P(>|Chi|)

1	8	10.336			
2	9	11.276	-1	-0.93971	0.3324

Model selection (II): the step() function in R

```
> step(first.glm, direction = "backward")
```

```
Start: AIC=59.53
```

```
Res ~ log(years) + Gov + Gov:log(years)
```

	Df	Deviance	AIC
--	----	----------	-----

- log(years):Gov	1	11.276	58.466
------------------	---	--------	--------

<none>		10.336	59.526
--------	--	--------	--------

```
Step: AIC=58.47
```

```
Res ~ log(years) + Gov
```

	Df	Deviance	AIC
--	----	----------	-----

- Gov	1	11.299	56.489
-------	---	--------	--------

<none>		11.276	58.466
--------	--	--------	--------

- log(years)	1	33.269	78.459
--------------	---	--------	--------

```
Step: AIC=56.49
```

```
Res ~ log(years)
```

	Df	Deviance	AIC
--	----	----------	-----

<none>		11.299	56.489
--------	--	--------	--------

-log(years)	1	33.436	76.626
-------------	---	--------	--------

```
-Call: glm(formula = Res ~ log(years), family = poisson)
```

```
Coefficients:
```

```
(Intercept) log(years)
```

```
0.3168 0.9654
```

```
Degrees of Freedom: 11 Total (i.e. Null); 10 Residual
```

```
Null Deviance: 33.44
```

```
Residual Deviance: 11.3
```

```
AIC: 56.49
```


Model 4 with the offset variable is not included here.

Example 3: smoking and coronary death

```
> CHD
  age smoke death  pop
1  40    s   32 52407
2  50    s  104 43248
3  60    s  206 28612
4  70    s  186 12663
5  80    s  102  5317
6  40   ns    2 18790
7  50   ns   12 10673
8  60   ns   28  5710
9  70   ns   28  2585
10 80   ns   31  1462
```

A study about tobacco consumption and coronary heart disease among British doctors.

Smoking and coronary death: the data



GLM with an offset variable: model formulation

Number of deaths per population size:

$$\frac{Y_i}{n_i} \sim \text{Poisson}(\mu_i)$$

$$Y_i \sim \text{Poisson}(n_i \times \mu_i)$$

$$g(\mu_i) = X\beta$$

Smoking and coronary death

- Is the death rate higher for smokers than non smokers ?
- If so, by how much ?
- Is there differential effects of age ?

GLM with an offset variable: model formulation

Number of deaths per
population size:

Linear predictor:

$$\eta = f(\text{age}, \text{smoking}, \text{pop.})$$

$$\frac{Y_i}{n_i} \sim \text{Poisson}(\mu_i)$$

$$Y_i \sim \text{Poisson}(n_i \times \mu_i)$$

$$g(\mu_i) = X\beta = \log(n_i) + \log(\mu_i)$$

Offset variable

$$\beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \dots$$

3 models

$$M_1 : \eta = \log(\text{pop.}) + \beta_1 \text{age} + \beta_2 \text{smoke}$$

$$M_2 : \eta = \log(\text{pop.}) + \beta_1 \text{age} + \beta_2 \text{smoke} + \beta_3 \text{age} \times \text{smoke}$$

$$M_3 : \eta = \log(\text{pop.}) + \beta_{11} \text{age} + \beta_{12} \text{age}^2 + \beta_2 \text{smoke} + \beta_3 \text{age} \times \text{smoke}$$

```
>fit.chd1<-glm(death ~ age + smoke+offset(log(pop)), poisson)
>fit.chd2<-glm(death ~ age + smoke+age:smoke+offset(log(pop)), poisson)
>age2<-age^2
>fit.chd3<-glm(death ~ age+age2+smoke+age:smoke+offset(log(pop)), poisson)
```


Model selection

```
> extractAIC(fit.chd1, k=2)
[1] 3.0000 130.2500
> extractAIC(fit.chd2, k=2)
[1] 4.0000 122.9614
> extractAIC(fit.chd3, k=2)
[1] 5.00000 66.70331
```

The model with quadratic age effect has the best goodness-of-fit.

R output (model 3)

```
> summary(fit.chd3)
```

Call:

```
glm(formula = death ~ age + age2 + smoke + age:smoke + offset(log(pop)),  
     family = poisson)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.970e+01	1.253e+00	-15.717	< 2e-16 ***
age	3.563e-01	3.632e-02	9.810	< 2e-16 ***
age2	-1.977e-03	2.737e-04	-7.223	5.08e-13 ***
smokes	2.364e+00	6.562e-01	3.602	0.000316 ***
age:smokes	-3.075e-02	9.704e-03	-3.169	0.001528 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


(Dispersion parameter for poisson family taken to be 1)

Null deviance: 935.0673 on 9 degrees of freedom
Residual deviance: 1.6354 on 5 degrees of freedom
AIC: 66.703

Number of Fisher Scoring iterations: 4

Data and predicted model

$$\hat{E}(Y_i) = n_i \times \exp(\mu_i) = n_i \times \exp(X\hat{\beta})$$



Chapter 12:

Beyond Poisson and binomial distributions: Models with different link functions and/or distributions

Lindsey: Chapter 4.

Example 1: Employment duration

- The employment duration of staff, age 25 to 44, recruited to the British post office in the first quarter of 1973 and classified in to two grades.

Month	Grade 1	Grade2	Month	Grade1	Grade2
1	22	30	13	0	1
2	18	28	14	0	0
3	19	31	15	0	0
4	13	14	16	1	1
5	5	10	17	1	1
6	6	6	18	1	0
7	3	5	19	3	2
8	2	2	20	1	0
9	2	3	21	1	3
10	1	0	22	0	1
11	0	0	23	0	1
12	1	1	24	0	0

The data in R

```
> employ <- read.table("C:.... /employ.txt",header=TRUE)
```

```
> attach(employ)
```


```
> employ
```

month grade duration


1	1	1	22
2	1	2	30
3	2	1	18
4	2	2	28
5	3	1	19
6	3	2	31
7	4	1	13
8	4	2	14
9	5	1	5
10	5	2	10

. . . .

Original sacle.



log transformation.



Model formulation

Models with normal error.

$$Y_t \sim N(\mu_t, \sigma^2)$$

$$\eta = \beta_0 + \beta_1 t + \beta_2 G$$

$$\mu = \beta_0 + \beta_1 t + \beta_2 G$$

identity link

$$\log(\mu) = \beta_0 + \beta_1 t + \beta_2 G$$

$$\mu = e^{\beta_0 + \beta_1 t + \beta_2 G}$$

Log link

$$\frac{1}{\mu} = \beta_0 + \beta_1 t + \beta_2 G$$

$$\mu = \frac{1}{\beta_0 + \beta_1 t + \beta_2 G}$$

inverse link

Models with normal error in R

```
m.normal.idt <- glm(duration ~ month + grade + month:grade,data = employ, family = gaussian(link = identity))
```

```
m.normal.inv <- glm(duration ~ month + grade + month:grade,data = employ, family = gaussian(link = inverse))
```

```
m.normal.log <- glm(duration ~ month + grade + month:grade,data = employ, family = gaussian(link = log))
```

```
m.normal.log1 <- glm(duration ~ month + grade ,data = employ,  
family = gaussian(link = log))
```

Model selection

```
> extractAIC(m.normal.idt, k=2)
[1] 4.0000 228.5094
> extractAIC(m.normal.inv, k=2)
[1] 4.0000 198.3163
> extractAIC(m.normal.log, k=2)
[1] 4.0000 173.8542
> extractAIC(m.normal.log1, k=2)
[1] 3.0000 171.8545
```

The model with log link has the smallest AIC value.

GLM with normal error and log link in R

```
> m.normal.log <- glm(duration ~ month + grade + month:grade, data = employ1, family = gaussian(link = log))
> summary(m.normal.log)
```

Call:

```
glm(formula = duration ~ month + grade + month:grade, family = gaussian(link = log),
    data = employ1)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-4.6676	-1.5332	-0.2005	0.8915	10.8676

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.4476971	0.1304802	26.423	< 2e-16 ***
month	-0.2709521	0.0444746	-6.092	1.08e-06 ***
grade2	0.3698490	0.1591805	2.323	0.0271 *
month:grade2	-0.0007872	0.0545143	-0.014	0.9886


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 8.219982)

Null deviance: 2771.88 on 33 degrees of freedom
Residual deviance: 246.59 on 30 degrees of freedom
AIC: 173.85

Number of Fisher Scoring iterations: 7

Data and fitted model



Other models

- Different distributions:
 - Normal.
 - Gamma.
 - ...

$$Y_t \sim H(\mu_t)$$

$$g(\mu_t) = \eta = \beta_0 + \beta_1 t + \beta_2 G$$

$$\mu = \beta_0 + \beta_1 t + \beta_2 G$$

identity link

$$\log(\mu) = \beta_0 + \beta_1 t + \beta_2 G$$

$$\mu = e^{\beta_0 + \beta_1 t + \beta_2 G}$$

Log link

$$\frac{1}{\mu} = \beta_0 + \beta_1 t + \beta_2 G$$

$$\mu = \frac{1}{\beta_0 + \beta_1 t + \beta_2 G}$$

inverse link

Model formulation (model 1)

```
m.normal.log <- glm(duration ~ month + grade + month:grade, data = employ1, family =  
gaussian(link = log))
```

$$\eta = \log(\mu_i) = \beta_0 + \beta_2 t_i + \beta_1 G_i + \beta_3 t_i \times G_i$$

A model with normal error and log link:

$$Y_i \sim N(\mu_i, \sigma^2)$$

$$g(\mu) = \eta, \quad \log(\mu) = \eta$$

R output : model 1

```
> m.normal.log <- glm(duration ~ month + grade + month:grade, data = employ, family = gaussian(link = log))
> summary(m.normal.log)
```

Call:

```
glm(formula = duration ~ month + grade + month:grade, family = gaussian(link = log),
    data = employ)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-4.6676	-1.5332	-0.2005	0.8915	10.8676

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.0778481	0.2764302	11.134	3.54e-12 ***
month	-0.2701650	0.0943705	-2.863	0.00759 **
grade	0.3698490	0.1591805	2.323	0.02712 *
month:grade	-0.0007872	0.0545143	-0.014	0.98857

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 8.219982)

Null deviance: 2771.88 on 33 degrees of freedom
Residual deviance: 246.59 on 30 degrees of freedom
AIC: 173.85

Model with gamma error and inverse link function: model formulation (model 2)

```
m.gamma.log <- glm(duration ~ month + grade + month:grade,  
  data = employ1,  
  family = Gamma(link = inverse))
```

$$\eta = 1/\mu_i = \beta_0 + \beta_2 t_i + \beta_1 G_i + \beta_3 t \times G_i$$

A model with gamma error and inverse:

$$Y_i \sim \text{Gamma}(\mu_i)$$

$$\frac{1}{\mu_i} = \eta_i$$

R output: model 2

```
> m.gamma.inv <- glm(duration ~ month + grade + month:grade, data = employ, family = Gamma(link = inverse))
> summary(m.gamma.inv )
```

Call:

```
glm(formula = duration ~ month + grade + month:grade, family = Gamma(link = inverse),
    data = employ)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.0122937	0.0396403	-0.310	0.75861
month	0.0492390	0.0154143	3.194	0.00329 **
grade	-0.0008403	0.0226236	-0.037	0.97062
month:grade	-0.0089011	0.0090337	-0.985	0.33235

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Gamma family taken to be 0.3661260)

Null deviance: 51.752 on 33 degrees of freedom

Residual deviance: 9.887 on 30 degrees of freedom

AIC: 147.27

Model with gamma error and inverse link function: model formulation (model 3)

Grade is not included in the model:

$$\eta = 1 / \mu_i = \beta_0 + \beta_2 t_i$$

A model with gamma error and inverse:

$$Y_i \sim \text{Gamma}(\mu_i)$$

$$\frac{1}{\mu_i} = \eta_i$$

R output: model 3

```
> m.gamma.inv1 <- glm(duration ~ month , data = employ, family = Gamma(link = inverse))
> summary(m.gamma.inv1)
```

Call:

```
glm(formula = duration ~ month, family = Gamma(link = inverse),
    data = employ)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-0.9163	-0.5284	-0.2795	0.2599	1.2685

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.014033	0.010821	-1.297	0.204
month	0.035353	0.004445	7.954	4.45e-09 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Gamma family taken to be 0.3762360)

Null deviance: 51.752 on 33 degrees of freedom
Residual deviance: 10.562 on 32 degrees of freedom
AIC: 145.63

Number of Fisher Scoring iterations: 6

Model selection

```
> AIC(m.normal.log)
```

```
[1] 173.8542
```

```
> AIC(m.gamma.inv)
```

```
[1] 147.2748
```

```
> AIC(m.gamma.inv1)
```


```
[1] 145.6314
```

```
> AIC(m.gamma.inv2)
```

```
[1] 145.3437
```

The model which give small AIC is the gamma model with inverse line and only duration in the model

Data and fitted models



Chapter 13:

Poisson regression and log linear models

Based on Dobson: Chapter 9

Log linear models

The general frame work of Poisson regression is given by:

$$Y_i \sim \text{Poisson}(\mu_i)$$

$$g(\mu_i) = \log(\mu_i) = X\beta$$

Example 1: melanoma

	Site			
Tumor Type	Head & neck	Trunk	Extremities	total
Hutchinson	22	2	10	34
Superficial melanoma	16	54	115	185
Nodular	19	33	73	125
Indeterminate	11	17	28	56
Total	68	106	226	400

Cross sectional study of patients with form of skin cancer.

Different sites & different tumor types.

Example 1

	Site			
Tumor Type	Head & neck	Trunk	Extremities	total
Hutchinson	22	2	10	34
Superficial melanoma	16	54	115	185
Nodular	19	33	73	125
Indeterminate	11	17	28	56
Total	68	106	226	400

Y_{ij} Frequency of subjects with tumor type i and site j .

$$Y_{ij} \sim \text{Poisson}(\mu_{ij})$$

n Sample size.

Example 1: independence ?

	Site			
Tumor Type	Head & neck	Trunk	Extremities	total
Hutchinson	22	2	10	34
Superficial melanoma	16	54	115	185
Nodular	19	33	73	125
Indeterminate	11	17	28	56
Total	68	106	226	400

$Y_{i.}$

$Y_{.j}$

Chi-squared statistic for independence:

$$X^2 = \frac{\sum_{ij} (Y_{ij} - e_{ij})^2}{e_{ij}}$$

$$e_{ij} = \frac{Y_{i.} \times Y_{.j}}{n}$$

$$e_{11} = \frac{34 \times 68}{400}$$

Example 1: independence ?

$$X^2 = \frac{\sum_{ij} (Y_{ij} - e_{ij})^2}{e_{ij}} = \frac{(22 - 5.78)^2}{5.78} + \dots + \frac{(28 - 31.64)^2}{31.64} = 65.8$$

Model formulation

$$Y_{ij} \sim \text{Poisson}(\mu_{ij})$$

$$g(\mu_{ij}) = \mu$$

The minimal
model

M1

$$g(\mu_{ij}) = \mu + \alpha_i + \beta_j$$

Independence
model

M2

$$g(\mu_{ij}) = \mu + \alpha_i + \beta_j + \alpha\beta_{ij}$$

Second ordered
interaction model

M3

Data in R

```
> melanoma
  Tumor Type y
1  Hu   HN 22
2  Hu Trunk 2
3  Hu Extrem 10
4  Su   HN 16
5  Su Trunk 45
6  Su Extrem 115
7  Nod   HN 19
8  Nod Trunk 33
9  Nod Extrem 73
10 Ind   HN 11
11 Ind Trunk 17
12 Ind Extrem 28
```

Models in R

```
> M1<-glm(y~1, family=poisson, data=melanoma)
```

$$g(\mu_{ij}) = \mu$$

```
> M2<-glm(y~Tumor+Type, family=poisson, data=melanoma)
```

$$g(\mu_{ij}) = \mu + \alpha_i + \beta_j$$

```
> M3<-glm(y~Tumor+Type+Tumor:Type,family=poisson,  
data=melanoma)
```

$$g(\mu_{ij}) = \mu + \alpha_i + \beta_j + \alpha\beta_{ij}$$

```
> AIC(M1)
```

```
[1] 348.8361
```

```
> AIC(M2)
```

```
[1] 121.5482
```

```
> AIC(M3)
```

```
[1] 82.9297
```

Likelihood ratio test

```
> anova.glm(M2,M3,test="Chisq")
```

Analysis of Deviance Table

Model 1: $y \sim \text{Tumor} + \text{Type}$

Model 2: $y \sim \text{Tumor} + \text{Type} + \text{Tumor}:\text{Type}$

	Resid. Df	Resid. Dev	Df	Deviance	P(> Chi)
1	6	50.618			
2	0	0.000	6	50.618	3.533e-09 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
>
```


Example 2: Death penalty

Death penalty verdicts for cases involving multiple murders in Florida between 1976 and 1987. This data is from an article that studied effects of racial characteristics on whether persons convicted of homicide received the death penalty. The 674 subjects classified is in to a 2x2x2 contingency table- two rows, two columns, and two layers.

		Victim's Race			
		White		Black	
		Defendant's Race		Defendant's Race	
		White	Black	White	Black
Death Penalty	Yes	53	11	0	4
	No	414	37	16	139
Percent Yes		11.3	22.9	0.0	2.8

Example 3: Death penalty

		Victim's Race			
		White		Black	
		Defendant's Race		Defendant's Race	
		White	Black	White	Black
Death Penalty	Yes	53	11	0	4
	No	414	37	16	139
Percent Yes		11.3	22.9	0.0	2.8

The response variable: death penalty verdicts

$$Y_{ijk} \sim \text{Poisson}(\mu_{ijk})$$

Example 3: Death penalty

		Victim's Race			
		White		Black	
		Defendant's Race		Defendant's Race	
		White	Black	White	Black
Death Penalty	Yes	53	11	0	4
	No	414	37	16	139

The mean structure

$$\log(\mu_{ijk}) = \mu + v_i + d_j + p_k + \text{interaction}$$

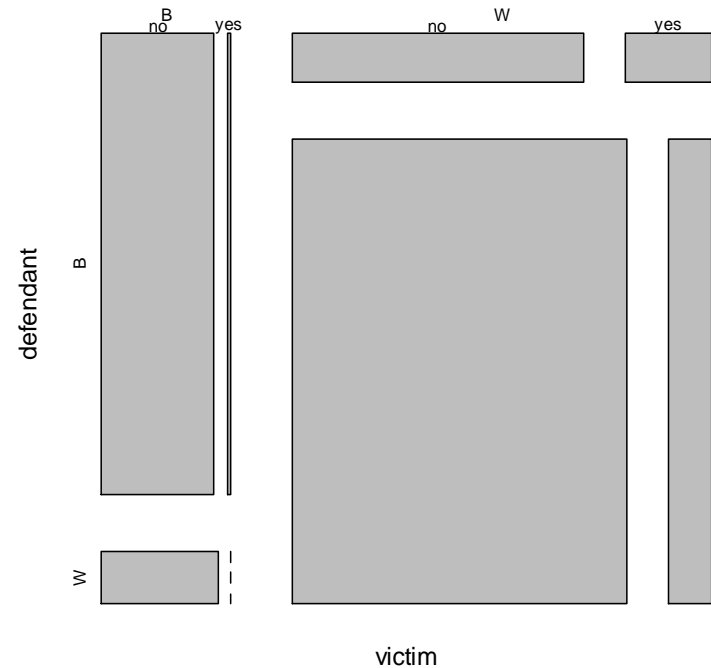
$$\log(\mu_{ijk}) = \mu + v_i + d_j + p_k + vd_{ij} + vp_{ik} + dp_{jk} + vdp_{ijk}$$

Exploring the Data in R

```
> deathpenalty <- data.frame(number = c(53,11,0,4,414,37,16,139),  
+ victim = c("W","W","B","B","W","W","B","B"),  
+ defendant = c("W","B","W","B","W","B","W","B"),  
+ death = rep(c("yes","no"),rep(4,2)))
```

```
plot(xtabs(number~ victim + defendant+ death, deathpenalty))
```

xtabs(number ~ victim + defendant + death, deathpenalty)



The model in R

$$\log(\mu_{ijk}) = \mu + v_i + d_j + p_k + vd_{ij} + vp_{ik} + dp_{jk} + vdp_{ijk}$$

```
M1<-glm(number~victim*defendant*death,  
        family=poisson,  
        data=deathpenalty)
```

Saturated model with Poisson family

```
> M1<-glm(number~victim*defendant*death, family=poisson, data=deathpenalty)
```

```
> summary(M1)
```

```
> Call:
```

```
Estimate Std. Error z value Pr(>|z|)
(Intercept)      4.934e+00  8.482e-02  58.177 < 2e-16 ***
victimW          -1.324e+00  1.850e-01  -7.155 8.38e-13 ***
defendantW       -2.162e+00  2.640e-01  -8.189 2.63e-16 ***
deathyes         -3.548e+00  5.071e-01  -6.996 2.63e-12 ***
victimW:defendantW  4.577e+00  3.149e-01  14.536 < 2e-16 ***
victimW:deathyes   2.335e+00  6.125e-01   3.813 0.000137 ***
defendantW:deathyes -2.153e+01  4.225e+04  -0.001 0.999593
victimW:defendantW:deathyes 2.068e+01  4.225e+04  0.00049 0.999609
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
(Dispersion parameter for poisson family taken to be 1)
```

```
Null deviance: 1.2251e+03 on 7 degrees of freedom
```

```
Residual deviance: 4.1224e-10 on 0 degrees of freedom
```

```
AIC: 54.04
```

```
Number of Fisher Scoring iterations: 20
```

Using the step function in R to identify the best model

`step(the three way interaction model)`

- A stepwise selection based on AIC.
- Starting point: the three way interaction model.

Step () function for model selection

> step(M1)

Start: **AIC=54.04**

number ~ victim * defendant * death

Df Deviance AIC

- victim:defendant:death 1 0.37984 **52.42**

<none> 0.00000 54.04

Step: **AIC=52.42**

number ~ victim + defendant + death + victim:defendant + victim:death +
defendant:death

Df Deviance AIC

<none> 0.38 52.42

- defendant:death 1 5.39 55.43

- victim:death 1 20.73 70.77

- victim:defendant 1 384.43 434.47

Starting
point (see
also slide 79).

Three way
interaction is
excluded.

Two way
interaction
are excluded
(one at the
time).

The best model includes all two way interactions:

$$\log(\mu_{ijk}) = \mu + v_i + d_j + p_k + vd_{ij} + vp_{ik} + dp_{jk}$$

Output of the Step() function

Call: glm(formula = number ~ victim + defendant + death + victim:defendant +
victim:death + defendant:death,
family = poisson, data = deathpenalty)


Coefficients:

(Intercept)	victimW	defendantW
4.9358	-1.3298	-2.1746
deathyes	victimW:defendantW	victimW:deathyes
-3.5961	4.5950	2.4044
defendantW:deathyes		
-0.8678		

Degrees of Freedom: 7 Total (i.e. Null); 1 Residual

Null Deviance: 1225

Residual Deviance: 0.3798 AIC: 52.42



The two way interaction model in R

$$\log(\mu_{ijk}) = \mu + v_i + d_j + p_k + vd_{ij} + vp_{ik} + dp_{jk}$$

```
M2<-glm(number~victim+defendant+death  
+victim:defendant  
+victim:death  
+defendant:death,  
family=poisson, data=deathpenalty)
```

Output in R

```
> summary(M2)
```

Call:

```
glm(formula = number ~ victim + defendant + death + victim:defendant +  
    victim:death + defendant:death, family = poisson, data = deathpenalty)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	4.93578	0.08471	58.265	< 2e-16 ***
victimW	-1.32980	0.18479	-7.196	6.19e-13 ***
defendantW	-2.17465	0.26377	-8.245	< 2e-16 ***
deathyes	-3.59610	0.50691	-7.094	1.30e-12 ***
victimW:defendantW	4.59497	0.31353	14.656	< 2e-16 ***
victimW:deathyes	2.40444	0.60061	4.003	6.25e-05 ***
defendantW:deathyes	-0.86780	0.36707	-2.364	0.0181 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 1225.07955 on 7 degrees of freedom

Residual deviance: 0.37984 on 1 degrees of freedom

AIC: 52.42

Example 3: Antibiotic prescription

Diagnosis of Respiratory Tract Infections Hueston and Stott (2000) report a study of clinicians' diagnoses of respiratory tract infections over a 14-month period. The aim was to determine whether a reduction in prescription of antibiotics to acute bronchitis patients was due to clinicians assigning an alternative diagnosis.

	Time period				
Diagnosis	1-3/96	4-6/96	7-9/96	10-12/96	1-2/97
Acute bronchitis	113	58	40	108	100
Acute sinusitis	99	37	23	50	32
URI	410	228	125	366	304
Pneumonia	60	43	30	56	45
Total	682	366	218	580	481

The data in R

```
> diag <- rep(c("bron", "sinus", "URI", "pneu"), 5)
> time <- rep(c("win96", "spr96", "sum96", "aut96", "spr97"), rep(4, 5))
> rt <- data.frame(diag = factor(diag, unique(diag)),
+ time = factor(time, unique(time)), count = c(113, 99, 410, 60, 58, 37, 228, 43, 40, 23,
+ 125, 30,
+ 108, 50, 366, 56, 100, 32, 304, 45))
```

Example 3: Antibiotic prescription

Diagnosis	Time period				
	1-3/96	4-6/96	7-9/96	10-12/96	1-2/97
Acute bronchitis	113	58	40	108	100
Acute sinusitis	99	37	23	50	32
URI	410	228	125	366	304
Pneumonia	60	43	30	56	45

Research question:
Diagnostic and time period are
independent ?

Data in R

```
> rt
  diag time count
1 bron win96 113
2 sinus win96 99
3 URI win96 410
4 pneu win96 60
5 bron spr96 58
6 sinus spr96 37
7 URI spr96 228
8 pneu spr96 43
9 bron sum96 40
10 sinus sum96 23
11 URI sum96 125
12 pneu sum96 30
13 bron aut96 108
14 sinus aut96 50
15 URI aut96 366
16 pneu aut96 56
17 bron spr97 100
18 sinus spr97 32
19 URI spr97 304
20 pneu spr97 45
```

Models formulation

$$Y_{ij} \sim \text{Poisson}(\mu_{ij})$$

$$g(\mu_{ij}) = \mu$$

Minimal model

$$g(\mu_{ij}) = \mu + D_i + T_j$$

Independence
model

$$g(\mu_{ij}) = \mu + D_i + T_j + DT_{ij}$$

Second ordered
interaction

Models in R

```
>
> M1 <- glm(count ~ 1, family=poisson, data=rt)
> M2 <- glm(count ~ diag+time, family=poisson, data=rt)
> M3 <- glm(count ~ diag+time+diag:time,family=poisson, data=rt)
>
>
> AIC(M1)
[1] 1915.304
> AIC(M2)
[1] 169.8704
> AIC(M3)
[1] 164.2791
```


Likelihood ratio test

```
> anova.glm(M2,M3,test="Chisq")
```

Analysis of Deviance Table

Model 1: count ~ diag + time

Model 2: count ~ diag + time + diag:time

	Resid. Df	Resid. Dev	Df	Deviance	P(> Chi)
1	12	29.591			
2	0	0.000	12	29.591	0.003216 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```
>
```

Independence model is rejected.

Output model 3

```
> summary(M3)
```

Coefficients:

```
      Estimate Std. Error z value Pr(>|z|)
(Intercept)   4.72739   0.09407  50.253 < 2e-16 ***
diagsinus    -0.13227   0.13766  -0.961 0.33664
diagURI       1.28877   0.10625  12.130 < 2e-16 ***
diagpneu     -0.63304   0.15974  -3.963 7.40e-05 ***
timespr96    -0.66694   0.16153  -4.129 3.64e-05 ***
timesum96     -1.03851   0.18398  -5.645 1.66e-08 ***
timeaut96    -0.04526   0.13457  -0.336 0.73664
timespr97    -0.12222   0.13729  -0.890 0.37336
diagsinus:timespr96 -0.31726  0.25143  -1.262 0.20702
diagURI:timespr96  0.08013  0.18143  0.442 0.65872
diagpneu:timespr96 0.33380  0.25693  1.299 0.19388
diagsinus:timesum96 -0.42112  0.29568  -1.424 0.15438
diagURI:timesum96  -0.14934  0.21045  -0.710 0.47795
diagpneu:timesum96 0.34536  0.28957  1.193 0.23300
diagsinus:timeaut96 -0.63784  0.21957  -2.905 0.00367 **
diagURI:timeaut96  -0.06827  0.15258  -0.447 0.65457
diagpneu:timeaut96 -0.02374  0.22942  -0.103 0.91760
diagsinus:timespr97 -1.00717  0.24536  -4.105 4.05e-05 ***
diagURI:timespr97  -0.17691  0.15677  -1.128 0.25913
diagpneu:timespr97 -0.16546  0.24029  -0.689 0.49107
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 1.7890e+03 on 19 degrees of freedom
Residual deviance: 1.5765e-14 on 0 degrees of freedom
AIC: 164.28

Number of Fisher Scoring iterations: 3

Chapter 14

Over dispersion

Lindsey: Chapter 3

Over dispersion parameter

- The general form of exponential family is defined as:

$$f(y) = \exp \left\{ \frac{y_i \theta_i - b(\theta_i)}{a_i(\phi)} + c(y_i, \phi) \right\}$$

- Where ϕ is the dispersion parameter.
- $a(\phi)$: scale parameter.

Example: normal distribution

$$f(y_i; \mu_i, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(y_i - \mu_i)^2}{2\sigma^2}}$$
$$= \exp\left\{\left[y_i\mu_i - \frac{\mu_i^2}{2}\right] \frac{1}{\sigma^2} - \frac{y_i^2}{2\sigma^2} - \frac{1}{2} \log(2\pi\sigma^2)\right\}$$



The mean and the variance are separated parameters.

$$\theta_i = \mu_i,$$

$$b(\theta_i) = \theta_i^2 / 2$$

$$a_i(\phi) = \sigma^2$$

$$c(y_i, \phi) = -[y_i^2 / \phi + \log(2\pi\phi)] / 2.$$

Example: Binomial distribution

$$Z_i = \begin{cases} 1 \\ 0 \end{cases} \quad \longrightarrow \quad Y_i = \sum_{i=1}^n Z_i \quad \longrightarrow \quad Y_i \sim B(n, \pi_i)$$

$$p(y_i | \theta) = \binom{n_i}{y_i} \theta^{y_i} (1 - \theta)^{n - y} =$$

$$\exp \left\{ y_i \log \left[\frac{\theta_i}{1 - \theta_i} \right] + n_i \log(1 - \theta_i) + \log \binom{n_i}{y_i} \right\}$$

The variance is a function of the mean.

$$a_i(\phi) = 1, \quad b(\theta_i) = \log(1 + \exp(\theta_i))$$

$$c(y) = \log \binom{n_i}{y_i}$$

$$E(y) = \mu = b'(\theta_i) = e^{\theta} (1 + \exp(\theta_i))^{-1}$$

$$\text{var}(y) = n\mu(1 - \mu)$$

Poisson distribution

$$Y_i \sim \text{Poisson}(\mu)$$

$$f(y_i, \theta_i) \frac{\theta_i^{y_i} e^{-\theta}}{y_i!}$$

$$a_i(\phi) = 1$$

$$b(\theta) = \exp(\theta)$$

$$c(y) = -\log(y!)$$

$$E(y) = \mu = b'(\theta) = \exp(\theta)$$

$$\text{var}(y) = \mu$$

$$E(y) = V(y).$$

Overdispersion

The binomial and Poisson distribution are a members of one parameter exponential family.

$$Y_t \sim \text{Poisson}(\mu_t)$$

$$E(Y_t) = V(Y_t) = \mu_t$$

$$Y_i \sim B(n, \pi)$$

$$E(Y_i) = n\pi$$

$$V(Y_i) = n\pi(1 - \pi)$$

$$a_i(\phi) = 1$$

Overdispersion

Often, we observed extra variability than expected:

$$V(Y_i) > E(Y_i)$$

$$V(Y_i) > n\pi(1-\pi)$$

$$V(Y_i) = \phi E(Y_i)$$

$$V(Y_i) = \phi n\pi(1-\pi)$$

$\phi = 1$  No problem with overdispersion

Estimating over dispersion

Formula

$$\phi = \frac{\chi_p^2}{n - p}$$

where

p = number of parametr in the model

n = number of observations

Example 1: Germination of seeds from Orobanche

o. aegyptiaco 75				o. aegyptiaco 73			
bean		cucumber		bean		cucumber	
germ.	total	germ.	total	germ.	total	germ.	total
10	39	5	6	8	16	3	12
23	62	53	74	10	30	22	41
23	<u>81</u>	55	72	8	28	15	30
26	51	32	51	23	45	32	51
17	39	46	79	0	<u>4</u>	3	7
		10	13				


Data in R

```
> fac<-read.table("C:...../seed.txt", header = TRUE)
> fact <- transform(fac, prop = germ/total, extr.seed = interaction(extract,seed))
```

```
fact
  seed extract germ total  prop  extr.seed
1 aeg75  bean  10   39 0.2564103  bean.aeg75
2 aeg75  bean  23   62 0.3709677  bean.aeg75
3 aeg75  bean  23   81 0.2839506  bean.aeg75
4 aeg75  bean  26   51 0.5098039  bean.aeg75
5 aeg75  bean  17   39 0.4358974  bean.aeg75
6 aeg75 cucumber  5    6 0.8333333 cucumber.aeg75
7 aeg75 cucumber 53   74 0.7162162 cucumber.aeg75
8 aeg75 cucumber 55   72 0.7638889 cucumber.aeg75
9 aeg75 cucumber 32   51 0.6274510 cucumber.aeg75
10 aeg75 cucumber 46   79 0.5822785 cucumber.aeg75
11 aeg75 cucumber 10   13 0.7692308 cucumber.aeg75
12 aeg73  bean   8   16 0.5000000  bean.aeg73
13 aeg73  bean  10   30 0.3333333  bean.aeg73
14 aeg73  bean   8   28 0.2857143  bean.aeg73
15 aeg73  bean  23   45 0.5111111  bean.aeg73
16 aeg73  bean   0    4 0.0000000  bean.aeg73
17 aeg73 cucumber  3   12 0.2500000 cucumber.aeg73
18 aeg73 cucumber 22   41 0.5365854 cucumber.aeg73
19 aeg73 cucumber 15   30 0.5000000 cucumber.aeg73
20 aeg73 cucumber 32   51 0.6274510 cucumber.aeg73
21 aeg73 cucumber  3    7 0.4285714 cucumber.aeg73
```

Exploring the data: Box-Plot

```
> plot(prop ~ extr.seed, data = fact, las = 1, ylab = "proportion")  
> points(prop ~ extr.seed, data = fact, pch = 16)  
> title("proportion of germinating seeds")
```



Model 1 formulation

- Binomial model

$$y_i \sim \text{Bin}(n_i, \pi_i)$$

$$g(\pi_i) = \eta, \quad \text{logit}(\pi_i) = \eta$$

$$\text{logit}(\pi_i) = \beta_0 + \beta_1 \text{ext} + \beta_2 \text{seed} + \beta_3 \text{ext} \times \text{seed}$$

Model 1: binomial family

```
> g <- glm(cbind(germ, total - germ) ~ extract + seed + extract:seed, family = binomial, data = fact)
> r.pears<-residuals(g, type="pearson")
> summary(g)
```

The assumption here is over dispersion parameter is $\Phi=1$

Model 1 output

Call:

```
glm(formula = cbind(germ, total - germ) ~ extract + seed + extract:seed,  
     family = binomial, data = fact)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.01617	-1.24398	0.05995	0.84695	2.12123

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.4122	0.1842	-2.238	0.0252 *
extractcucumber	0.5401	0.2498	2.162	0.0306 *
seedaeg75	-0.1459	0.2232	-0.654	0.5132
extractcucumber:seedaeg75	0.7781	0.3064	2.539	0.0111 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 98.719 on 20 degrees of freedom

Residual deviance: 33.278 on 17 degrees of freedom

AIC: 117.87

Number of Fisher Scoring iterations: 4

Estimating over dispersion in R

```
> X2 <- sum(residuals(g, type = "pearson")^2)
> X2
[1] 31.65114
> phi <- X2/g$df.residual
> phi
[1] 1.861832
> phi <- g$deviance/g$df.residual
> phi
[1] 1.957517
```

As we can see from the R output, the over dispersion parameter is greater than 1 and hence this is an indication of the presence of over dispersion

Confidence interval for model 1

```
> g <- glm(cbind(germ, total - germ) ~ extract + seed + extract:seed, family =  
  binomial, data = fact)
```

```
> confint(g)
```

Waiting for profiling to be done...

	2.5 %	97.5 %
(Intercept)	-0.77866159	-0.05469853
extractcucumber	0.05275108	1.03329024
seedaeg75	-0.58184899	0.29428969
extractcucumber:seedaeg75	0.17619697	1.37823747

Taking into account overdispersion in R

```
glm(model, family = quasibinomial,...)
```

Output quasi-binomial model 2

```
> g.over <- glm(cbind(germ, total - germ) ~ extract + seed + extract:seed, family = quasibinomial, data = fact)
> summary(g.over)
```

Call:

```
glm(formula = cbind(germ, total - germ) ~ extract + seed + extract:seed,
     family = quasibinomial, data = fact)
```

Deviance Residuals:

```
      Min       1Q   Median       3Q      Max
-2.01617 -1.24398  0.05995  0.84695  2.12123
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.4122	0.2513	-1.640	0.1193
extractcucumber	0.5401	0.3409	1.584	0.1315
seedae75	-0.1459	0.3045	-0.479	0.6379
extractcucumber:seedae75	0.7781	0.4181	1.861	0.0801 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for quasibinomial family taken to be 1.861832)

Null deviance: 98.719 on 20 degrees of freedom
Residual deviance: 33.278 on 17 degrees of freedom
AIC: NA

Number of Fisher Scoring iterations: 4

```
> summary(g.over)$dispersion
[1] 1.861832
```

Confidence Interval model2

```
> library(MASS)
> fact$prop <- with(fact, germ/total)
> g.over.alt <- glm(prop ~ extract + seed + extract:seed,
+ weights = total, family = quasibinomial, data = fact)
> confint(g.over.alt)
```

Waiting for profiling to be done...

	2.5 %	97.5 %
(Intercept)	-0.91562380	0.07457178
extractcucumber	-0.12415989	1.21529201
seedaeg75	-0.74043242	0.45663323
extractcucumber:seedaeg75	-0.04413964	1.59702376

Estimates and CI for proportion with and without over dispersion

	$\hat{\pi}$	2.5%	97.5%
$\phi = 1$	0.36	0.31	0.42
$\hat{\phi} = 1.86$	0.36	0.28	0.45

To calculate this you need to install the R packages:

```
library(doBy)  
library(R2HTML)
```

Example 2: Habitat preferences of lizards

- A study consists of two lizards type: Grohami and Opalinus.
- Response: number of sites (from the total) occupied by Grahami lizards.
- Covariates:
 1. Height of the site (H).
 2. Diameter (D).
 3. Sun condition of the site (S, sun/ shade).
 4. Time of the day (T).

Habitat preferences of lizards

> habitat

	G	Total	S	D	H	T
1	20	22	S1	D1	H1	Early
2	8	9	S1	D1	H1	Mid
3	4	8	S1	D1	H1	Late
4	13	13	S1	D1	H2	Early
5	8	8	S1	D1	H2	Mid
6	12	12	S1	D1	H2	Late
7	8	11	S1	D2	H1	Early
8	4	5	S1	D2	H1	Mid
9	5	8	S1	D2	H1	Late
10	6	6	S1	D2	H2	Early
11	0	0	S1	D2	H2	Mid
12	1	2	S1	D2	H2	Late
13	34	45	S2	D1	H1	Early
14	69	89	S2	D1	H1	Mid
15	18	28	S2	D1	H1	Late
16	31	36	S2	D1	H2	Early
17	55	59	S2	D1	H2	Mid
18	13	16	S2	D1	H2	Late
19	17	32	S2	D2	H1	Early
20	60	92	S2	D2	H1	Mid
21	8	16	S2	D2	H1	Late
22	12	13	S2	D2	H2	Early
23	21	26	S2	D2	H2	Mid
24	4	8	S2	D2	H2	Late

S: sun conditions sun / shade).

D: diameter (<2 / > 2).

H: hight (< 5 / > 5).

T: time of day (early/ mid day/late).

Habitat preferences of lizards: model formulation

$$y_{ijkl} \sim B(n_{ijkl}, \pi_{ijkl})$$

Total ample size.



Number of sites occupied by Grahmi lizards.

π_{ijkl} = The probability that a site is occupied by Grahmi lizards.

$$g(\pi_{ijkl}) = \mu + \alpha_i + \beta_j + \gamma_k + \delta_l + \dots$$

Habitat preferences of lizards: model formulation in R

Main effects model in R

$$g(\pi_{ijkl}) = \mu + \alpha_i + \beta_j + \gamma_k + \delta_l$$

```
> f1<-glm((G/Total)~H+D+S+T,family="binomial",data=habitat)
```

R output

```
> summary(f1)
```

Call:

```
glm(formula = (G/Total) ~ H + D + S + T, family = "binomial",  
     data = habitat)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-0.50878	-0.11019	0.02009	0.26466	0.52322

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	2.0618	1.4060	1.466	0.143
HH2	1.0631	1.1222	0.947	0.343
DD2	-0.8798	1.0841	-0.812	0.417
SS2	-0.6415	1.0884	-0.589	0.556
TLate	-1.2054	1.2761	-0.945	0.345
TMid	0.0587	1.4590	0.040	0.968

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 4.6730 on 22 degrees of freedom
Residual deviance: 1.5417 on 17 degrees of freedom
(1 observation deleted due to missingness)
AIC: 28.658

Interpretation

Coefficients:


	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	2.0618	1.4060	1.466	0.143
HH2	1.0631	1.1222	0.947	0.343
DD2	-0.8798	1.0841	-0.812	0.417
SS2	-0.6415	1.0884	-0.589	0.556
TLate	-1.2054	1.2761	-0.945	0.345
TMid	0.0587	1.4590	0.040	0.968

All the parameters estimates are not significant.

diagnostic

```
>r.pearson<-resid(f1, type="pearson")  
> par(mfrow=c(1,1))  
> qqnorm(r.pearson)  
> abline(0,1)
```


$$r_i^P = \frac{y_i - \hat{\mu}_i}{\sqrt{V(\hat{\mu}_i)}} \sim N(0,1)$$



diagnostic

```
> mean(r.pearson)
[1] 0.01454735
> var(r.pearson)
[1] 0.05871372
> r.s<-(r.pearson-mean(r.pearson))/
  sqrt((var(r.pearson)))
> qqnorm(r.s)
> abline(0,1)
```

The variance of
pearson residual is
much smaller than 1



Taking into account overdispersion

```
glm((G/Total)~H+D+S+T,family="binomial",data=habitat)
```



```
glm((G/Total)~H+D+S+T,family=quasibinomial,data= habitat)
```


Taking into account overdispersion

```
> f1.over <- glm((G/Total)~H+D+S+T,family=quasibinomial,data= habitat)
> summary(f1.over)
```

Call:

```
glm(formula = (G/Total) ~ H + D + S + T, family = quasibinomial,
    data = habitat)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-0.50878	-0.11019	0.02009	0.26466	0.52322

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.0618	0.3883	5.310	5.76e-05 ***
HH2	1.0631	0.3099	3.430	0.00319 **
DD2	-0.8798	0.2994	-2.939	0.00918 **
SS2	-0.6415	0.3006	-2.134	0.04768 *
TLate	-1.2054	0.3524	-3.420	0.00326 **
TMid	0.0587	0.4029	0.146	0.88588

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for quasibinomial family taken to be 0.07626879)

Null deviance: 4.6730 on 22 degrees of freedom
Residual deviance: 1.5417 on 17 degrees of freedom
(1 observation deleted due to missingness)
AIC: NA


Number of Fisher Scoring iterations: 5

Diagnostic

```
> mean(r.pearson)
[1] 0.01454735
> var(r.pearson)
[1] 0.05871372
> r.s<-(r.pearson-mean(r.pearson))/
  sqrt((var(r.pearson)))
> qqnorm(r.s)
> abline(0,1)
```

The variance of
pearson residual is
much smaller than 1

```
> summary(f1.over)$dispersion
[1] 0.07626879
```



Example 3: ship damage

- The ships data from the MASS package concern a type of damage caused by waves to the forward section of cargo-carrying vessels.
- The variables are
 - incidents number of damage incidents
 - service aggregate months of service
 - period period of operation : 1960-74, 75-79
 - year year of construction: 1960-64, 65-69, 70-74, 75-79
 - type type: "A" to "E"
- Here it makes sense to model the expected number of incidents per aggregate months of service.

Data in R

The data is available in the R the MASS library as:

```
> library(MASS)
> data(ships)
> ships2 <- subset(ships, service > 0)
> ships2$year <- as.factor(ships2$year)
➤ ships2$period <- as.factor(ships2$period)
```

```
> ships
  type year period service incidents
```

```
1  A  60   60   127      0
2  A  60   75    63      0
3  A  65   60  1095      3
4  A  65   75  1095      4
5    A  70   60  1512      6
.    .    .    .    .
37  E  70   60  1157      5
38  E  70   75  2161     12
39  E  75   60     0      0
40  E  75   75   542      1
```

Mean structure and model formulation in R

Model formulation:

$$Y_{ijk} \sim \text{Poisson}(\mu_{ijk})$$

$$g(\mu_{ijk}) = \mu + \text{Type}_i + \text{Year}_j + \text{Period}_k + \underbrace{\log(\text{service})}_{\text{offset}}$$

Model formulation in R:

```
> glm1 <- glm(formula = incidents ~ type + year + period,  
+             family = poisson(link = "log"), data = ships2,  
+             offset = log(service))
```

$\phi = 1$

R output

```
> glm1 <- glm(formula = incidents ~ type + year + period,  
+ family = poisson(link = "log"), data = ships2,  
+ offset = log(service))  
> summary(glm1)  
Call:  
glm(formula = incidents ~ type + year + period, family = poisson(link = "log"),  
data = ships2, offset = log(service))  
Deviance Residuals:
```

```
Min      1Q  Median      3Q      Max  
-1.6768 -0.8293 -0.4370  0.5058  2.7912
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-6.40590	0.21744	-29.460	< 2e-16 ***
typeB	-0.54334	0.17759	-3.060	0.00222 **
typeC	-0.68740	0.32904	-2.089	0.03670 *
typeD	-0.07596	0.29058	-0.261	0.79377
typeE	0.32558	0.23588	1.380	0.16750
year65	0.69714	0.14964	4.659	3.18e-06 ***
year70	0.81843	0.16977	4.821	1.43e-06 ***
year75	0.45343	0.23317	1.945	0.05182 .
period75	0.38447	0.11827	3.251	0.00115 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for poisson family taken to be 1)

Null deviance: 146.328 on 33 degrees of freedom

Residual deviance: 38.695 on 25 degrees of freedom

AIC: 154.56

Number of Fisher Scoring iterations: 5

Model 1: quasi-poisson log linear

```
> glm2 <- update(glm1, family = quasipoisson(link = "log"))
```

$$\phi \neq 1$$

```
> summary(glm2)
```

> Call:

```
glm(formula = incidents ~ type + year + period, family = quasipoisson(link = "log"),  
    data = ships2, offset = log(service))
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-6.40590	0.28276	-22.655	< 2e-16 ***
typeB	-0.54334	0.23094	-2.353	0.02681 *
typeC	-0.68740	0.42789	-1.607	0.12072
typeD	-0.07596	0.37787	-0.201	0.84230
typeE	0.32558	0.30674	1.061	0.29864
year65	0.69714	0.19459	3.583	0.00143 **
year70	0.81843	0.22077	3.707	0.00105 **
year75	0.45343	0.30321	1.495	0.14733
period75	0.38447	0.15380	2.500	0.01935 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for quasipoisson family taken to be 1.691028)

Null deviance: 146.328 on 33 degrees of freedom

Residual deviance: 38.695 on 25 degrees of freedom

AIC: NA

Number of Fisher Scoring iterations: 5

Standard errors are
changed since $\Phi > 1$.

Estimation of over dispersion parameter

```
> X2 <- sum(residuals(glm2, type = "pearson")^2)
```

```
> X2
```

```
[1] 42.27525
```

```
> phi <- X2/glm2$df.residual
```

```
> phi
```

```
[1] 1.69101
```

```
> phi <- g$deviance/glm2$df.residual
```

```
> phi
```

```
[1] 1.331111
```


CI for the parameters

```
> confint(glm2)
```

	2.5 %	97.5 %
(Intercept)	-6.9789252	-5.86832189
typeB	-0.9793495	-0.07040017
typeC	-1.6043341	0.09972077
typeD	-0.8628299	0.63543426
typeE	-0.2880598	0.92322467
year65	0.3217115	1.08674462
year70	0.3882167	1.25564720
year75	-0.1562814	1.03712256
period75	0.0841713	0.68792187

Summary

- ANOVA table without and with assuming over dispersion

Effect	DF	$\phi = 1$	$\hat{\phi} = 1.67$
		$P(> \chi^2)$	$P(>F)$
Type	4	2.63E-11	2.29E-04
Period	1	1.1E-03	1.888E-02
Year	3	5.038e-09	5.777E-04